

Explainable Artificial Intelligence (AI): Local Explanation with DALEX (SHAP-like explanations)

Lumumba Wandera Victor

2024-11-30

EXPLAINABLE AI FOR MALARIA MODELING: Local Explanation with DALEX

INSTRUCTOR: LUMUMBA WANDERA VICTOR

AMMnet ML GROUP

30TH NOVEMBER, 2024

Definition:

Explainable AI (XAI) refers to a set of tools and techniques that make the behavior of machine learning models more interpretable and understandable to humans.

Goals

The goal of XAI is to clarify how and why AI models make certain predictions ensuring that both experts and non-experts can trust and effectively use these insights in real-world applications.

In malaria modeling, for instance, XAI can explain why an AI model predicts high malaria severity in revealing which symptoms (like vomiting, jaundice, or fever) are influencing that prediction.

By increasing model interpretability, XAI helps bridge the gap between complex AI technologies and actionable insights for policymakers, healthcare workers, and researchers, allowing them to use AI insights more confidently in planning and intervention.

Types

- SHapleY Additive Model exPlanations (SHAP)
- Local Interpretable Model-Agnostic Explanations (LIME)
- Local Explanation with DALEX (SHAP-like explanations)

These methods help identify which features (or variables) contribute most to a model's predictions.

Load necessary packages

```
#install.packages("caret") # If not already installed
#install.packages("iml")
#install.packages("lme")
#install.packages('tictoc')
#install.packages("ggplot2")
#install.packages("dplyr")
```

Load LIBRARIES

```
library(caret)
library(iml)
library(lme)
library(tictoc)
library(ggplot2)
library(dplyr)
```

Load the Malaria dataset

Source: <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7093799/>

```
mdata = read.csv("Malaria-Data.csv", header = TRUE)
attach(mdata)
```

```
dim(mdata)
```

```
[1] 337 18
```

```
head(mdata)
```

	age	sex	fever	cold	rigor	fatigue	headace	bitter_tongue	vomitting	diarrhea
1	3	1	1	1	0	1	1	1	0	1
2	3	0	1	1	1	1	1	1	0	1
3	3	0	1	1	1	1	1	0	0	1
4	4	1	1	1	0	1	0	0	0	0
5	4	0	1	1	1	0	1	0	0	0
6	4	1	0	0	0	1	1	0	0	1

	Convulsion	Anemia	jundice	cocacola_urine	hypoglycemia	prostration
1	1	0	1	1	1	0
2	0	0	0	1	1	0
3	1	0	0	1	1	0
4	0	0	1	0	1	0
5	1	0	1	1	1	0
6	0	1	0	0	0	0

	hyperpyrexia	severe_malaria
1	0	0
2	0	0
3	0	1
4	1	0
5	0	0
6	0	1

View the Column Names in the Data set

```
names(mdata)
```

```
[1] "age"      "sex"      "fever"    "cold"
[5] "rigor"    "fatigue"  "headace"  "bitter_tongue"
[9] "vomitting" "diarrhea" "Convulsion" "Anemia"
[13] "jundice"  "cocacola_urine" "hypoglycemia" "prostration"
[17] "hyperpyrexia" "severe_malaria"
```

View the Structure of the Dataset

```
str(mdata)
```

```
'data.frame':  337 obs. of  18 variables:
 $ age      : int  3 3 3 4 4 4 4 5 5 8 ...
 $ sex      : int  1 0 0 1 0 1 0 1 0 0 ...
 $ fever    : int  1 1 1 1 1 0 1 1 1 1 ...
 $ cold     : int  1 1 1 1 1 0 1 0 0 1 ...
 $ rigor    : int  0 1 1 0 1 0 1 1 1 1 ...
 $ fatigue  : int  1 1 1 1 0 1 1 1 1 0 ...
 $ headace  : int  1 1 1 0 1 1 0 0 1 1 ...
 $ bitter_tongue : int  1 1 0 0 0 0 0 0 1 1 ...
 $ vomitting : int  0 0 0 0 0 0 0 1 0 0 ...
 $ diarrhea : int  1 1 1 0 0 1 0 1 0 1 ...
 $ Convulsion : int  1 0 1 0 1 0 0 0 1 0 ...
 $ Anemia   : int  0 0 0 0 0 1 0 1 0 0 ...
 $ jundice  : int  1 0 0 1 1 0 0 1 0 1 ...
 $ cocacola_urine: int  1 1 1 0 1 0 0 0 0 0 ...
 $ hypoglycemia : int  1 1 1 1 1 0 0 0 1 1 ...
 $ prostration : int  0 0 0 0 0 0 0 0 0 0 ...
 $ hyperpyrexia : int  0 0 0 1 0 0 0 0 0 0 ...
 $ severe_malaria: int  0 0 1 0 0 1 0 0 0 0 ...
```

Descriptive Statistics

```
summary(mdata) ###Descriptive Statistics
```

age	sex	fever	cold
Min. : 3.00	Min. :0.0000	Min. :0.0000	Min. :0.0000
1st Qu.:19.00	1st Qu.:0.0000	1st Qu.:1.0000	1st Qu.:0.0000
Median :29.00	Median :1.0000	Median :1.0000	Median :1.0000
Mean :30.35	Mean :0.5341	Mean :0.7507	Mean :0.5668
3rd Qu.:38.00	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000
Max. :77.00	Max. :1.0000	Max. :1.0000	Max. :1.0000
rigor	fatigue	headace	bitter_tongue
Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.0000
1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000

Median :0.0000	Median :0.0000	Median :1.0000	Median :0.0000
Mean :0.3412	Mean :0.4837	Mean :0.7003	Mean :0.4036
3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000
Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.0000
vomitting	diarrhea	Convulsion	Anemia
Min. :0.00000	Min. :0.0000	Min. :0.0000	Min. :0.0000
1st Qu.:0.00000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000
Median :0.00000	Median :0.0000	Median :0.0000	Median :0.0000
Mean :0.07418	Mean :0.3383	Mean :0.3442	Mean :0.3501
3rd Qu.:0.00000	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000
Max. :1.00000	Max. :1.0000	Max. :1.0000	Max. :1.0000
jundice	cocacola_urine	hypoglycemia	prostration
Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.0000
1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:1.0000	1st Qu.:0.0000
Median :1.0000	Median :1.0000	Median :1.0000	Median :0.0000
Mean :0.6588	Mean :0.5401	Mean :0.8576	Mean :0.2196
3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:0.0000
Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.0000
hyperpyrexia	severe_malaria		
Min. :0.0000	Min. :0.0000		
1st Qu.:0.0000	1st Qu.:0.0000		
Median :0.0000	Median :0.0000		
Mean :0.1395	Mean :0.3442		
3rd Qu.:0.0000	3rd Qu.:1.0000		
Max. :1.0000	Max. :1.0000		

```
library(psych)
describe(mdata)###Descriptive Statistics
```

	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew
age	1	337	30.35	14.72	29	29.22	14.83	3	77	74	0.75
sex	2	337	0.53	0.50	1	0.54	0.00	0	1	1	-0.14
fever	3	337	0.75	0.43	1	0.81	0.00	0	1	1	-1.15
cold	4	337	0.57	0.50	1	0.58	0.00	0	1	1	-0.27
rigor	5	337	0.34	0.47	0	0.30	0.00	0	1	1	0.67
fatigue	6	337	0.48	0.50	0	0.48	0.00	0	1	1	0.07
headache	7	337	0.70	0.46	1	0.75	0.00	0	1	1	-0.87
bitter_tongue	8	337	0.40	0.49	0	0.38	0.00	0	1	1	0.39
vomitting	9	337	0.07	0.26	0	0.00	0.00	0	1	1	3.24
diarrhea	10	337	0.34	0.47	0	0.30	0.00	0	1	1	0.68
Convulsion	11	337	0.34	0.48	0	0.31	0.00	0	1	1	0.65
Anemia	12	337	0.35	0.48	0	0.31	0.00	0	1	1	0.63
jundice	13	337	0.66	0.47	1	0.70	0.00	0	1	1	-0.67
cocacola_urine	14	337	0.54	0.50	1	0.55	0.00	0	1	1	-0.16
hypoglycemia	15	337	0.86	0.35	1	0.94	0.00	0	1	1	-2.04
prostration	16	337	0.22	0.41	0	0.15	0.00	0	1	1	1.35
hyperpyrexia	17	337	0.14	0.35	0	0.05	0.00	0	1	1	2.07
severe_malaria	18	337	0.34	0.48	0	0.31	0.00	0	1	1	0.65
	kurtosis	se									
age		0.49	0.80								
sex		-1.99	0.03								
fever		-0.67	0.02								
cold		-1.93	0.03								
rigor		-1.56	0.03								

fatigue	-2.00	0.03
headace	-1.25	0.02
bitter_tongue	-1.85	0.03
vomitting	8.49	0.01
diarrhea	-1.54	0.03
Convulsion	-1.58	0.03
Anemia	-1.61	0.03
jundice	-1.56	0.03
cocacola_urine	-1.98	0.03
hypoglycemia	2.16	0.02
prostration	-0.18	0.02
hyperpyrexia	2.30	0.02
severe_malaria	-1.58	0.03

Check the Number of Missing Values

```
sum(is.na(mdata))###Check for missing data
```

```
[1] 0
```

Exclude Age from the dataset

```
mdata=mdata[,-1] ##Exclude Age
names(mdata)
```


```
[1] "sex"          "fever"          "cold"           "rigor"
[5] "fatigue"      "headace"        "bitter_tongue"  "vomitting"
[9] "diarrhea"     "Convulsion"     "Anemia"         "jundice"
[13] "cocacola_urine" "hypoglycemia"   "prostration"    "hyperpyrexia"
[17] "severe_malaria"
```

Rename the classes of the Target variable and plot it to determine imbalance

```
mdata$severe_malaria <- factor(mdata$severe_malaria,
                                levels = c(0,1),
                                labels = c('Not Severe', 'Severe'))
```

Perform Featureplot to see the data distribution at a glance

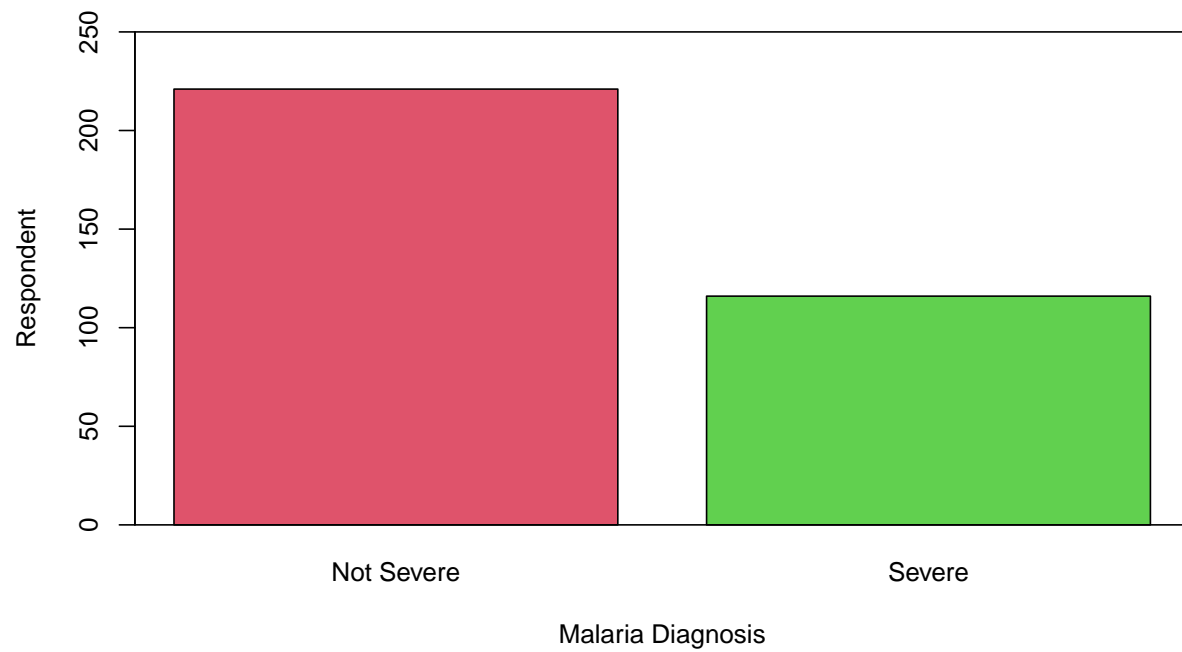
```
#png("Features Plot.png", height = 800, width = 1000)
featurePlot(x = mdata[, -which(names(mdata) == "severe_malaria")], # Predictors
            y = mdata$severe_malaria, # Target variable
            plot = "box", # Type of plot (e.g., "box", "density",
            #strip = strip.custom(strip.names = TRUE), # Add strip labels
            scales = list(x = list(relation = "free"), # Scales for x-axis
            y = list(relation = "free")), main = "Visual Comparison of Feature Distribution")
```



```
#dev.off()
```

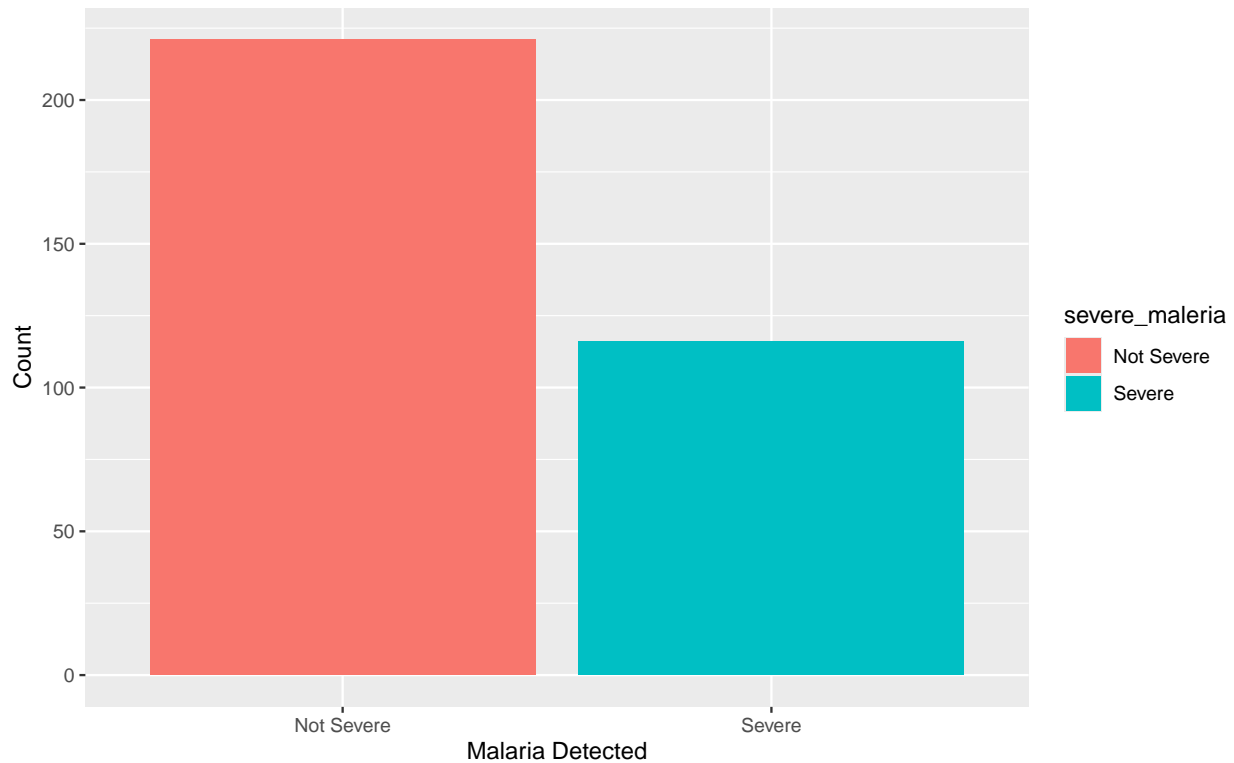
Plot Target Variable

```
plot(factor(severe_malaria), names= c('Not Severe', 'Severe'),
     col=c(2,3),
     ylim=c(0, 250),
     ylab='Respondent',
     xlab='Malaria Diagnosis')
box()
```



#Or use ggplot

```
ggplot(mdata, aes(x = factor(severe_malaria), fill = severe_malaria)) +  
  geom_bar() +  
  labs(x = "Malaria Detected",  
       y = "Count")
```



VIEW THE AVAILABLE MODELS IN CARET

```
models = getModelInfo()
names(models)
```

[1] "ada"	"AdaBag"	"AdaBoost.M1"
[4] "adaboost"	"amdai"	"ANFIS"
[7] "avNNet"	"awnb"	"awtan"
[10] "bag"	"bagEarth"	"bagEarthGCV"
[13] "bagFDA"	"bagFDAGCV"	"bam"
[16] "bartMachine"	"bayesglm"	"binda"
[19] "blackboost"	"blasso"	"blassoAveraged"
[22] "bridge"	"brnn"	"BstLm"
[25] "bstSm"	"bstTree"	"C5.0"
[28] "C5.0Cost"	"C5.0Rules"	"C5.0Tree"
[31] "cforest"	"chaid"	"CSimca"
[34] "ctree"	"ctree2"	"cubist"
[37] "dda"	"deepboost"	"DENFIS"
[40] "dnn"	"dwdLinear"	"dwdPoly"
[43] "dwdRadial"	"earth"	"elm"
[46] "enet"	"evtrees"	"extraTrees"
[49] "fda"	"FH.GBML"	"FIR.DM"
[52] "foba"	"FRBCS.CHI"	"FRBCS.W"
[55] "FS.HGD"	"gam"	"gamboost"
[58] "gamLoess"	"gamSpline"	"gaussprLinear"
[61] "gaussprPoly"	"gaussprRadial"	"gbm_h2o"

[64]	"gbm"	"gcvEarth"	"GFS.FR.MOGUL"
[67]	"GFS.LT.RS"	"GFS.THRIFT"	"glm.nb"
[70]	"glm"	"glmboost"	"glmnet_h2o"
[73]	"glmnet"	"glmStepAIC"	"gpls"
[76]	"hda"	"hdda"	"hdrda"
[79]	"HYFIS"	"icr"	"J48"
[82]	"JRip"	"kernelpls"	"kknk"
[85]	"knn"	"krlsPoly"	"krlsRadial"
[88]	"lars"	"lars2"	"lasso"
[91]	"lda"	"lda2"	"leapBackward"
[94]	"leapForward"	"leapSeq"	"Linda"
[97]	"lm"	"lmStepAIC"	"LMT"
[100]	"loclda"	"logicBag"	"LogitBoost"
[103]	"logreg"	"lssvmLinear"	"lssvmPoly"
[106]	"lssvmRadial"	"lvq"	"M5"
[109]	"M5Rules"	"manb"	"mda"
[112]	"Mlda"	"mlp"	"mlpKerasDecay"
[115]	"mlpKerasDecayCost"	"mlpKerasDropout"	"mlpKerasDropoutCost"
[118]	"mlpML"	"mlpSGD"	"mlpWeightDecay"
[121]	"mlpWeightDecayML"	"monmlp"	"msaenet"
[124]	"multinom"	"mxnet"	"mxnetAdam"
[127]	"naive_bayes"	"nb"	"nbDiscrete"
[130]	"nbSearch"	"neuralnet"	"nnet"
[133]	"nnls"	"nodeHarvest"	"null"
[136]	"OneR"	"ordinalNet"	"ordinalRF"
[139]	"ORFlog"	"ORFpls"	"ORFridge"
[142]	"ORFsvm"	"ownn"	"pam"
[145]	"parRF"	"PART"	"partDSA"
[148]	"pcaNNet"	"pcr"	"pda"
[151]	"pda2"	"penalized"	"PenalizedLDA"
[154]	"plr"	"pls"	"plsRglm"
[157]	"polr"	"ppr"	"pre"
[160]	"PRIM"	"protoclass"	"qda"
[163]	"QdaCov"	"qrf"	"qrnn"
[166]	"randomGLM"	"ranger"	"rbf"
[169]	"rbfDDA"	"Rborist"	"rda"
[172]	"regLogistic"	"relaxo"	"rf"
[175]	"rFerns"	"RFlda"	"rfRules"
[178]	"ridge"	"rllda"	"rlm"
[181]	"rmlda"	"rocc"	"rotationForest"
[184]	"rotationForestCp"	"rpart"	"rpart1SE"
[187]	"rpart2"	"rpartCost"	"rpartScore"
[190]	"rqlasso"	"rqnc"	"RRF"
[193]	"RRFglobal"	"rrlda"	"RSimca"
[196]	"rvmlLinear"	"rvmlPoly"	"rvmlRadial"
[199]	"SBC"	"sda"	"sdwd"
[202]	"simpls"	"SLAVE"	"slda"
[205]	"smda"	"snn"	"sparseLDA"
[208]	"spikeslab"	"splls"	"stepLDA"
[211]	"stepQDA"	"superpc"	"svmBoundrangeString"
[214]	"svmExpoString"	"svmLinear"	"svmLinear2"
[217]	"svmLinear3"	"svmLinearWeights"	"svmLinearWeights2"
[220]	"svmPoly"	"svmRadial"	"svmRadialCost"
[223]	"svmRadialSigma"	"svmRadialWeights"	"svmSpectrumString"

[226] "tan"	"tanSearch"	"treebag"
[229] "vbmpRadial"	"vglmAdjCat"	"vglmContRatio"
[232] "vglmCumulative"	"widekernelpls"	"WM"
[235] "wsrf"	"xgbDART"	"xgbLinear"
[238] "xgbTree"	"xyf"	

TODAY we are going to train the following machine learning models:

- LR
- SVM
- RANDOM FOREST
- NAIVE BAYES
- KNN
- LDA
- NNET/mlp
- LVQ
- Bagging
- Boosting
- DT

STEPS

1. Data Preparation and Preprocessing, Cleaning, Feature Engineering, Visualization, Data Splitting, etc
2. Define the Training Control- Set up cross validation
3. Train the Models- Select the ML models you want to train
4. Evaluate your model using test data
5. Tune the hyperparameters and Resample the data (optional)
6. Implement XAI

DATA PARTITION FOR MACHINE LEARNING

caret can also be used for data partition

```
set.seed(123)
trainIndex <- createDataPartition(mdata$severe_malaria, p = 0.7, list = FALSE)
train <- mdata[trainIndex, ]
test <- mdata[-trainIndex, ]
dim(train)
```

```
[1] 237  17
```

```
dim(test)
```

```
[1] 100  17
```

Set seed for reproducibility

```
set.seed(123)
```

Define control for training

```
#control1 <- trainControl(method = "cv", number = 10)
control1 <- trainControl(method = "repeatedcv",
                          number = 10, repeats = 5,
                          sampling = 'smote',
                          search = 'random') ## For tuning
```

The R code above defines a `trainControl` object named `control1`, used to configure the training process for machine learning models within the `caret` package. The `method = "repeatedcv"` specifies that repeated k-fold cross-validation will be applied, ensuring robust model performance evaluation. Specifically, `number = 10` sets the number of folds to 10, and `repeats = 5` means this cross-validation will be repeated five times, reducing the variability in performance metrics. The `sampling = 'smote'` parameter implements Synthetic Minority Over-sampling Technique (SMOTE) during training, a technique that addresses class imbalance by generating synthetic examples in the minority class. Lastly, `search = 'random'` indicates that hyperparameter tuning will be performed using random search rather than a grid search, making it more efficient by sampling a random subset of the hyperparameter space. This setup improves the model's accuracy and generalizability, particularly for imbalanced datasets, by preventing overfitting and ensuring reliable hyperparameter selection.

1. Train the Logistic Regression Model

```
tic()
lrModel <- train(severe_malaria ~ .,
                 data = train,
                 method = "glm",
                 trControl = control1)
toc()
```

1.61 sec elapsed

Predict on the test set

```
lrpred <- predict(lrModel, newdata = test)
```

Evaluate with Confusion Matrix

```
lr.cM <- confusionMatrix(lrpred, test$severe_malaria,
                         positive = "Severe",
                         mode = "everything")
print(lr.cM)
```

Confusion Matrix and Statistics

Prediction	Reference	
	Not Severe	Severe
Not Severe	38	15
Severe	28	19

Accuracy : 0.57
95% CI : (0.4671, 0.6686)
No Information Rate : 0.66
P-Value [Acc > NIR] : 0.97603

Kappa : 0.1232

Mcnemar's Test P-Value : 0.06725

Sensitivity : 0.5588
Specificity : 0.5758
Pos Pred Value : 0.4043
Neg Pred Value : 0.7170
Precision : 0.4043
Recall : 0.5588
F1 : 0.4691
Prevalence : 0.3400
Detection Rate : 0.1900
Detection Prevalence : 0.4700
Balanced Accuracy : 0.5673

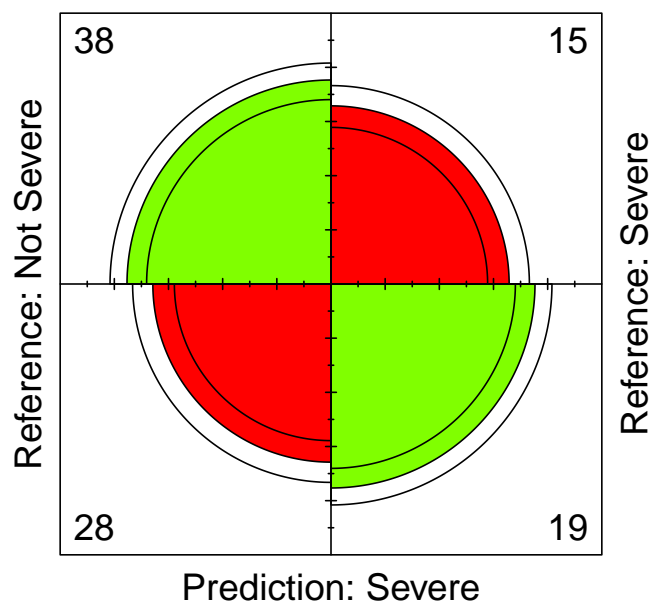
'Positive' Class : Severe

Plotting confusion matrix


```
fourfoldplot(lr.cm$table,  
             col = rainbow(4),  
             main = "LR Confusion Matrix")
```

LR Confusion Matrix

Prediction: Not Severe



```
plot(varImp(lrModel, scale = TRUE))
```



SHAP (SHapley Additive exPlanations)

The Shapley value helps explain how much each feature contributes to the prediction made by a machine learning model. It provides a way to fairly distribute the “credit” for the model’s output across all input features. By visualizing the SHAP plot, you can understand not only which features are important, but also how specific feature values that are driving predictions for individual cases.

Set seed for reproducibility

```
set.seed(456)
```

Assuming lrModel is already trained : Convert the caret model to a Predictor object, separating the target variable

```
predictorlr <- Predictor$new(lrModel,
                             data = train[, -which(names(train) == "severe_malaria")],
                             y = train$severe_malaria)
```

Select a single instance from the test set to explain Replace ‘1’ with the index of any other instance if desired

```
x_interest <- test[1, -which(names(test) == "severe_malaria")]
```

Compute SHAP values for the specific instance

```
shapleylr <- Shapley$new(predictorlr, x_interest = x_interest)
```

View the SHAP Values

```
shapleylr$results
```


	feature	class	phi	phi.var	feature.value
1	sex	Not Severe	0.00	0.00000000	sex=0
2	fever	Not Severe	0.00	0.00000000	fever=1
3	cold	Not Severe	-0.04	0.03878788	cold=1
4	rigor	Not Severe	0.05	0.04797980	rigor=1
5	fatigue	Not Severe	-0.12	0.10666667	fatigue=1
6	headace	Not Severe	-0.16	0.13575758	headace=1
7	bitter_tongue	Not Severe	0.08	0.07434343	bitter_tongue=1
8	vomitting	Not Severe	0.00	0.00000000	vomitting=0
9	diarrhea	Not Severe	-0.13	0.11424242	diarrhea=1
10	Convulsion	Not Severe	-0.17	0.14252525	Convulsion=0
11	Anemia	Not Severe	0.03	0.02939394	Anemia=0
12	jundice	Not Severe	0.08	0.07434343	jundice=0
13	cocacola_urine	Not Severe	-0.01	0.01000000	cocacola_urine=1
14	hypoglycemia	Not Severe	-0.06	0.05696970	hypoglycemia=1
15	prostration	Not Severe	-0.07	0.06575758	prostration=0
16	hyperpyrexia	Not Severe	0.00	0.00000000	hyperpyrexia=0

17	sex	Severe	0.00	0.00000000	sex=0
18	fever	Severe	0.00	0.00000000	fever=1
19	cold	Severe	0.04	0.03878788	cold=1
20	rigor	Severe	-0.05	0.04797980	rigor=1
21	fatigue	Severe	0.12	0.10666667	fatigue=1
22	headace	Severe	0.16	0.13575758	headace=1
23	bitter_tongue	Severe	-0.08	0.07434343	bitter_tongue=1
24	vomitting	Severe	0.00	0.00000000	vomitting=0
25	diarrhea	Severe	0.13	0.11424242	diarrhea=1
26	Convulsion	Severe	0.17	0.14252525	Convulsion=0
27	Anemia	Severe	-0.03	0.02939394	Anemia=0
28	jundice	Severe	-0.08	0.07434343	jundice=0
29	cocacola_urine	Severe	0.01	0.01000000	cocacola_urine=1
30	hypoglycemia	Severe	0.06	0.05696970	hypoglycemia=1
31	prostration	Severe	0.07	0.06575758	prostration=0
32	hyperpyrexia	Severe	0.00	0.00000000	hyperpyrexia=0

```
#View(shapleylr$results)
```

Plot the SHAP values for this instance

```
shapleylr$plot() +  
  ggtitle("SHAP Values for a Single Instance in Logistic Regression Model")
```



INTERPRETATION

Each feature has its own SHAP value, calculated in the context of all other features. The direction and length of the bar indicate the magnitude and impact on the prediction. Rightward (positive): Indicates the feature is pushing the model prediction towards a positive class (e.g., “Severe” if that is the positive label).

Local Explanation with DALEX (SHAP-like explanations)

```
#install.packages("DALEX")
library(DALEX)

# Create the explainer object
explainer <- explain(lrModel,
                     data = test[, -ncol(test)], # Exclude the outcome column
                     y = as.numeric(as.character(test$severe_malaria)),
                     label = "Local Explanation with DALEX for Logistic Regression")
```


Preparation of a new explainer is initiated

```
-> model label      : Local Explanation with DALEX for Logistic Regression
-> data             : 100 rows 16 cols
-> target variable  : 100 values
-> predict function : yhat.train will be used ( default )
-> predicted values : No value for predict function target column. ( default )
-> model_info       : package caret , ver. 6.0.94 , task classification ( default )
-> predicted values : numerical, min = 0.1097885 , mean = 0.4777402 , max = 0.7690111
-> residual function : difference between y and yhat ( default )
-> residuals        : numerical, min = NA , mean = NA , max = NA
A new explainer has been created!
```

```
# Select an instance to explain (e.g., first row in test data set)
instance <- test[1, -ncol(test)] # Exclude the outcome column for prediction

# Generate explanations for the instance
local_explanation <- predict_parts(explainer, new_observation = instance)

# Plot local explanation
plot(local_explanation)
```

Overview

The graph presents a local explanation of a logistic regression model using the DALEX package. It visualizes how different predictor variables contribute to the prediction for a specific instance. The prediction is represented by the bar on the right, and the contributions of each variable are shown as horizontal bars.

Breakdown of Contributions

- **Intercept:** This baseline value represents the model's prediction when all predictor variables are zero or absent. In this case, the intercept is 0.478.
- **Predictor Variables:** Each predictor variable's contribution is shown as a bar. The color indicates the direction of the contribution:
 - **Green:** Positive contribution, meaning the variable increases the prediction.
 - **Red:** Negative contribution, meaning the variable decreases the prediction.

The length of the bar represents the magnitude of the contribution.

Interpretation of the Specific Variables

Diarrhea = 1: Having diarrhea positively contributes to the prediction, with a value of +0.076. **Fatigue = 1:** Fatigue also has a positive impact, contributing +0.065. **Bitter_tongue = 1:** A bitter tongue negatively contributes to the prediction, with a value of -0.061. **Headache = 1:** Having a headache positively contributes +0.055. **Jaundice = 0:** Not having jaundice negatively contributes -0.049. **Convulsion = 0:** Not having convulsions positively contributes +0.05. **Rigor = 1:** Rigor negatively contributes -0.032. **Hypoglycemia**

= 1: Hypoglycemia positively contributes +0.036. Cocacola_urine = 1: Having coca-cola colored urine positively contributes +0.024. Cold = 1: Having a cold positively contributes +0.013. All other factors: The remaining factors not explicitly shown contribute a small positive value of +0.001.

Overall Prediction

Summing up all the contributions (intercept + predictor variables), we arrive at the final prediction of 0.656. This value represents the probability of a certain outcome, as logistic regression models typically output probabilities.

This graph provides a valuable tool for understanding how a logistic regression model arrives at a specific prediction. It highlights the relative importance of different predictor variables and their impact on the final outcome. However, it is essential to consider the limitations and interpret the results in conjunction with other model evaluation metrics.

Additional Explanation (More simpler explanation)

The plot explains how different health symptoms contribute to a prediction made by a logistic regression model. Let's break it down step by step in simple terms. The model predicts severe malaria occurrence. The prediction value shown at the bottom is 0.656 (or 65.6%), meaning the model is moderately confident about this outcome. The model starts with a base value, the "intercept," which is 0.478. This represents the starting point for the prediction, assuming no additional information about symptoms.

Adding symptoms:

The model adjusts the base value based on the presence or absence of various symptoms.

For example, diarrhea = 1 (having diarrhea) adds 0.076 to the base, increasing the prediction. Similarly, fatigue = 1 adds 0.065, and rigor = 1 (shivering) adds 0.036.

Subtracting symptoms:

Some symptoms decrease the prediction value.

For instance, bitter_tongue = 1 (a bitter taste in the mouth) subtracts 0.061, and jaundice = 0 (absence of jaundice) subtracts 0.049.

The final prediction:

By adding and subtracting these values step by step, the model arrives at the final prediction of 0.656.

What does this mean

The plot shows how much each symptom influenced the prediction. Positive bars (green) pushed the prediction higher, while negative bars (red) pulled it down. This can help doctors or experts understand what factors were most important in the decision. For example, diarrhea and fatigue played a significant role in increasing the prediction, while a bitter tongue reduced it.

RANDOM FOREST

2. Train the Random Forest Classifier

```
tic()
rfModel <- train(factor(severe_malaria) ~ .,
                  data = train,
                  method = "rf", trControl = control1)
toc()
```

27.51 sec elapsed

Predict on the test set

```
rfpred <- predict(rfModel, newdata = test)
```

Evaluate with Confusion Matrix

```
rf.cm <- confusionMatrix(rfpred,
                          as.factor(test$severe_malaria),
                          positive = "Severe",
                          mode = "everything")
print(rf.cm)
```

Confusion Matrix and Statistics

	Reference	
Prediction	Not Severe	Severe
Not Severe	47	26
Severe	19	8

Accuracy : 0.55
95% CI : (0.4473, 0.6497)
No Information Rate : 0.66
P-Value [Acc > NIR] : 0.9915

Kappa : -0.0553

Mcnemar's Test P-Value : 0.3711


Sensitivity : 0.2353
Specificity : 0.7121
Pos Pred Value : 0.2963
Neg Pred Value : 0.6438
Precision : 0.2963
Recall : 0.2353
F1 : 0.2623
Prevalence : 0.3400

Detection Rate : 0.0800
Detection Prevalence : 0.2700
Balanced Accuracy : 0.4737


'Positive' Class : Severe

Plotting confusion matrix

```
fourfoldplot(rf.cM$table, col = rainbow(4), main = "RF Confusion Matrix")
```



```
plot(varImp(rfModel, scale = TRUE))
```



SHAP (SHapley Additive exPlanations)

The Shapley value helps explain how much each feature contributes to the prediction made by a machine learning model. It provides a way to fairly distribute the “credit” for the model’s output across all input features. By visualizing the SHAP plot, you can understand not only which features are important, but also how specific feature values that are driving predictions for individual cases.

Set seed for reproducibility

```
set.seed(456)
```

Assuming lrModel is already trained : Convert the caret model to a Predictor object, separating the target variable

```
predictorrf <- Predictor$new(rfModel,
                             data = train[, -which(names(train) == "severe_malaria")],
                             y = train$severe_malaria)
```

Select a single instance from the test set to explain Replace ‘1’ with the index of any other instance if desired


```
x_interest <- test[1, -which(names(test) == "severe_malaria")]
```

Compute SHAP values for the specific instance

```
shapleyrf <- Shapley$new(predictorrf, x_interest = x_interest)
```

Plot the SHAP values for this instance

```
shapleyrf$plot() + ggtitle("SHAP Values for a Single Instance in Random Forest Model")
```



Leftward (negative): Indicates the feature is pushing the model prediction towards a negative class (e.g., “Not Severe”). Larger absolute SHAP values mean a feature has a stronger influence on the prediction. Smaller SHAP values (close to zero) indicate that a feature has minimal influence on the model’s output for that instance

TRY FOR OTHER MODELS

NAIVE BAYES

3. Train the Naive Bayes Classifier

```
tic()
nbModel <- train(severe_malaria ~ .,
  data = train,
  method = "nb", trControl = control1)
toc()
```

6.17 sec elapsed

Predict on the test set

```
nbpred <- predict(nbModel, newdata = test)
```

Evaluate with Confusion Matrix

```
nb.cM <- confusionMatrix(nbpred,  
                          as.factor(test$severe_malaria),  
                          positive = "Severe",  
                          mode = "everything")  
print(nb.cM)
```

Confusion Matrix and Statistics

	Reference	
Prediction	Not Severe	Severe
Not Severe	35	11
Severe	31	23

Accuracy : 0.58
95% CI : (0.4771, 0.678)
No Information Rate : 0.66
P-Value [Acc > NIR] : 0.96195

Kappa : 0.181

Mcnemar's Test P-Value : 0.00337

Sensitivity : 0.6765
Specificity : 0.5303
Pos Pred Value : 0.4259
Neg Pred Value : 0.7609
Precision : 0.4259
Recall : 0.6765
F1 : 0.5227
Prevalence : 0.3400
Detection Rate : 0.2300
Detection Prevalence : 0.5400
Balanced Accuracy : 0.6034


'Positive' Class : Severe

Plotting confusion matrix


```
fourfoldplot(nb.cM$table, col = rainbow(4),  
             main = "NB Confusion Matrix")
```

NB Confusion Matrix

Prediction: Not Severe



```
plot(varImp(nbModel, scale = TRUE))
```



SHAP (SHapley Additive exPlanations)

The Shapley value helps explain how much each feature contributes to the prediction made by a machine learning model. It provides a way to fairly distribute the “credit” for the model’s output across all input features. By visualizing the SHAP plot, you can understand not only which features are important, but also how specific feature values that are driving predictions for individual cases.

Set seed for reproducibility

```
set.seed(456)
```

Assuming lrModel is already trained : Convert the caret model to a Predictor object, separating the target variable

```
predictornb <- Predictor$new(nbModel,  
                             data = train[, -which(names(train) == "severe_malaria")],  
                             y = train$severe_malaria)
```

Select a single instance from the test set to explain. Replace ‘1’ with the index of any other instance if desired

```
x_interest <- test[1, -which(names(test) == "severe_malaria")]
```


Compute SHAP values for the specific instance

```
shapleynb <- Shapley$new(predictornb, x_interest = x_interest)
```

Plot the SHAP values for this instance

```
shapleynb$plot() +  
  ggtitle("SHAP Values for a Single Instance in Naive Bayes Model")
```

SHAP Values for a Single Instance in Naive Bayes Model



Leftward (negative): Indicates the feature is pushing the model prediction towards a negative class (e.g., “Not Severe”). Larger absolute SHAP values mean a feature has a stronger influence on the prediction. Smaller SHAP values (close to zero) indicate that a feature has minimal influence on the model’s output for that instance

DECISION TREE

3. Train the Decision Tree Classifier

```
tic()
DTModel <- train(factor(severe_malaria) ~ .,
                  data = train,

                  method = "rpart", trControl = control11)
toc()
```

1.97 sec elapsed

Predict on the test set

```
DTpred <- predict(DTModel, newdata = test)
```

Evaluate with Confusion Matrix

```
DT.cM <- confusionMatrix(DTpred,  
                          as.factor(test$severe_malaria),  
                          positive = "Severe",  
                          mode = "everything")  
print(DT.cM)
```

Confusion Matrix and Statistics

	Reference	
Prediction	Not Severe	Severe
Not Severe	50	20
Severe	16	14

Accuracy : 0.64
95% CI : (0.5379, 0.7336)
No Information Rate : 0.66
P-Value [Acc > NIR] : 0.7039

Kappa : 0.1743

Mcnemar's Test P-Value : 0.6171


Sensitivity : 0.4118
Specificity : 0.7576
Pos Pred Value : 0.4667
Neg Pred Value : 0.7143
Precision : 0.4667
Recall : 0.4118
F1 : 0.4375
Prevalence : 0.3400
Detection Rate : 0.1400
Detection Prevalence : 0.3000
Balanced Accuracy : 0.5847

'Positive' Class : Severe


Plotting confusion matrix

```
fourfoldplot(DT.cM$table, col = rainbow(4),  
             main = "Decision Tree Confusion Matrix")
```

Decision Tree Confusion Matrix



```
plot(varImp(DTModel, scale = TRUE))
```



SHAP (SHapley Additive exPlanations)

The Shapley value helps explain how much each feature contributes to the prediction made by a machine learning model. It provides a way to fairly distribute the “credit” for the model’s output across all input features. By visualizing the SHAP plot, you can understand not only which features are important, but also how specific feature values that are driving predictions for individual cases.

Set seed for reproducibility

```
set.seed(456)
```

Assuming lrModel is already trained : Convert the caret model to a Predictor object, separating the target variable

```
predictorDT <- Predictor$new(DTModel,  
                             data = train[, -which(names(train) == "severe_malaria")],  
                             y = train$severe_malaria)
```

Select a single instance from the test set to explain Replace ‘1’ with the index of any other instance if desired

```
x_interest <- test[1, -which(names(test) == "severe_malaria")]
```


Compute SHAP values for the specific instance

```
shapleyDT <- Shapley$new(predictorDT, x_interest = x_interest)
```

Plot the SHAP values for this instance

```
shapleyDT$plot() +  
  ggtitle("SHAP Values for a Single Instance in Decision Tree Model")
```

SHAP Values for a Single Instance in Decision Tree Model



Leftward (negative): Indicates the feature is pushing the model prediction towards a negative class (e.g., “Not Severe”). Larger absolute SHAP values mean a feature has a stronger influence on the prediction. Smaller SHAP values (close to zero) indicate that a feature has minimal influence on the model’s output for that instance

KNN

3. Train the K-Nearest Neighbors Classifier

```
tic()
knnModel <- train(factor(severe_malaria) ~ .,
                  data = train,
                  method = "knn", trControl = control1)
toc()
```

2.35 sec elapsed

Predict on the test set

```
knnpred <- predict(knnModel, newdata = test)
```

Evaluate with Confusion Matrix

```
knn.cM <- confusionMatrix(knnpred,
                          as.factor(test$severe_malaria),
                          positive = "Severe",
                          mode = "everything")
print(knn.cM)
```

Confusion Matrix and Statistics

	Reference	
Prediction	Not Severe	Severe
Not Severe	20	10
Severe	46	24

Accuracy : 0.44
95% CI : (0.3408, 0.5428)
No Information Rate : 0.66
P-Value [Acc > NIR] : 1

Kappa : 0.0071

Mcnemar's Test P-Value : 2.91e-06


Sensitivity : 0.7059
Specificity : 0.3030
Pos Pred Value : 0.3429
Neg Pred Value : 0.6667
Precision : 0.3429
Recall : 0.7059
F1 : 0.4615
Prevalence : 0.3400
Detection Rate : 0.2400
Detection Prevalence : 0.7000
Balanced Accuracy : 0.5045

'Positive' Class : Severe


Plotting confusion matrix

```
fourfoldplot(knn.cM$table, col = rainbow(4),
             main = "Decision Tree Confusion Matrix")
```

Decision Tree Confusion Matrix



```
plot(varImp(knnModel, scale = TRUE))
```



SHAP (SHapley Additive exPlanations)

The Shapley value helps explain how much each feature contributes to the prediction made by a machine learning model. It provides a way to fairly distribute the “credit” for the model’s output across all input features. By visualizing the SHAP plot, you can understand not only which features are important, but also how specific feature values that are driving predictions for individual cases.

Set seed for reproducibility

```
set.seed(456)
```

Assuming lrModel is already trained : Convert the caret model to a Predictor object, separating the target variable

```
predictorknn <- Predictor$new(knnModel,  
                             data = train[, -which(names(train) == "severe_malaria")],  
                             y = train$severe_malaria)
```

Select a single instance from the test set to explain. Replace ‘1’ with the index of any other instance if desired

```
x_interest <- test[1, -which(names(test) == "severe_malaria")]
```


Compute SHAP values for the specific instance

```
shapleyknn <- Shapley$new(predictorknn, x_interest = x_interest)
```

Plot the SHAP values for this instance

```
shapleyknn$plot() +  
  ggtitle("SHAP Values for a Single Instance in k-NN Model")
```

SHAP Values for a Single Instance in k-NN Model



Leftward (negative): Indicates the feature is pushing the model prediction towards a negative class (e.g., “Not Severe”). Larger absolute SHAP values mean a feature has a stronger influence on the prediction. Smaller SHAP values (close to zero) indicate that a feature has minimal influence on the model’s output for that instance

SVM

3. Train the Support Vector Machines Classifier

```
tic()
svmModel <- train(factor(severe_malaria) ~ .,
  data = train,
  method = "svmRadial", trControl = control1)
toc()
```

7.07 sec elapsed

Predict on the test set

```
svmpred <- predict(svmModel, newdata = test)
```

Evaluate with Confusion Matrix

```
svm.cM <- confusionMatrix(svmpred,
                          as.factor(test$severe_malaria),
                          positive = "Severe",
                          mode = "everything")
print(svm.cM)
```

Confusion Matrix and Statistics

	Reference	
Prediction	Not Severe	Severe
Not Severe	47	23
Severe	19	11

Accuracy : 0.58
95% CI : (0.4771, 0.678)
No Information Rate : 0.66
P-Value [Acc > NIR] : 0.9620

Kappa : 0.0367

Mcnemar's Test P-Value : 0.6434

Sensitivity : 0.3235
Specificity : 0.7121
Pos Pred Value : 0.3667
Neg Pred Value : 0.6714
Precision : 0.3667
Recall : 0.3235
F1 : 0.3438
Prevalence : 0.3400
Detection Rate : 0.1100
Detection Prevalence : 0.3000
Balanced Accuracy : 0.5178

'Positive' Class : Severe

Plotting confusion matrix

```
fourfoldplot(svm.cM$table, col = rainbow(4), main = "Decision Tree Confusion Matrix") plot(varImp(svmModel,
scale = TRUE))
```

SHAP (SHapley Additive exPlanations)

The Shapley value helps explain how much each feature contributes to the prediction made by a machine learning model. It provides a way to fairly distribute the “credit” for the model’s output across all input features. By visualizing the SHAP plot, you can understand not only which features are important, but also how specific feature values that are driving predictions for individual cases.

Set seed for reproducibility

```
set.seed(456)
```

Assuming lrModel is already trained : Convert the caret model to a Predictor object, separating the target variable

```
predictorsvm <- Predictor$new(svmModel,  
                              data = train[, -which(names(train) == "severe_malaria")],  
                              y = train$severe_malaria)
```


Select a single instance from the test set to explain. Replace '1' with the index of any other instance if desired
desired_x_interest <- test[1, -which(names(test) == "severe_malaria")]

Compute SHAP values for the specific instance

```
shapleysvm <- Shapley$new(predictorsvm, x_interest = desired_x_interest)
```

Plot the SHAP values for this instance

```
shapleysvm$plot() +  
  ggtitle("SHAP Values for a Single Instance in Support Vector Machines Model")
```



Leftward (negative): Indicates the feature is pushing the model prediction towards a negative class (e.g., “Not Severe”). Larger absolute SHAP values mean a feature has a stronger influence on the prediction. Smaller SHAP values (close to zero) indicate that a feature has minimal influence on the model’s output for that instance

TREE BAG MODEL

3. Train the Tree Bag Classifier

```
tic()
TbagModel <- train(factor(severe_malaria) ~ .,
                   data = train,
                   method = "treebag", trControl = control1)
toc()
```

10.5 sec elapsed

Predict on the test set

```
Tbagpred <- predict(TbagModel, newdata = test)
```

Evaluate with Confusion Matrix

```
Tbag.cM <- confusionMatrix(Tbagpred,
                           as.factor(test$severe_malaria),
                           positive = "Severe",
                           mode = "everything")
print(Tbag.cM)
```

Confusion Matrix and Statistics

	Reference	
Prediction	Not Severe	Severe
Not Severe	48	25
Severe	18	9

Accuracy : 0.57
95% CI : (0.4671, 0.6686)
No Information Rate : 0.66
P-Value [Acc > NIR] : 0.9760

Kappa : -0.0084

Mcnemar’s Test P-Value : 0.3602

Sensitivity : 0.2647
Specificity : 0.7273

```

Pos Pred Value : 0.3333
Neg Pred Value : 0.6575
Precision : 0.3333
Recall : 0.2647
F1 : 0.2951
Prevalence : 0.3400
Detection Rate : 0.0900
Detection Prevalence : 0.2700
Balanced Accuracy : 0.4960

'Positive' Class : Severe

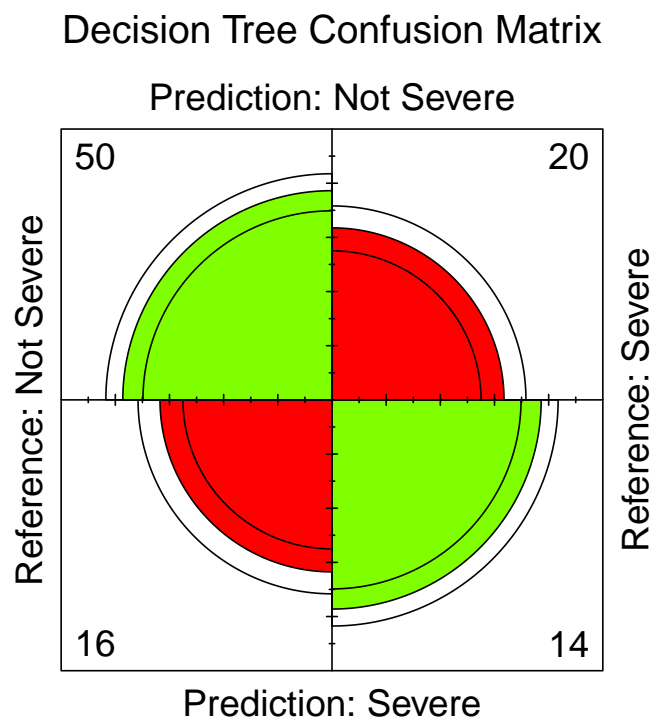
```

Plotting confusion matrix

```

fourfoldplot(DT.cM$table, col = rainbow(4),
  main = "Decision Tree Confusion Matrix")


```



```

plot(varImp(DTModel, scale = TRUE))

```



SHAP (SHapley Additive exPlanations)

The Shapley value helps explain how much each feature contributes to the prediction made by a machine learning model. It provides a way to fairly distribute the “credit” for the model’s output across all input features. By visualizing the SHAP plot, you can understand not only which features are important, but also how specific feature values that are driving predictions for individual cases.

Set seed for reproducibility

```
set.seed(456)
```

Assuming lrModel is already trained : Convert the caret model to a Predictor object, separating the target variable

```
predictorTbag <- Predictor$new(TbagModel,
  data = train[, -which(names(train) == "severe_malaria")],
  y = train$severe_malaria)
```

Select a single instance from the test set to explain. Replace ‘1’ with the index of any other instance if desired

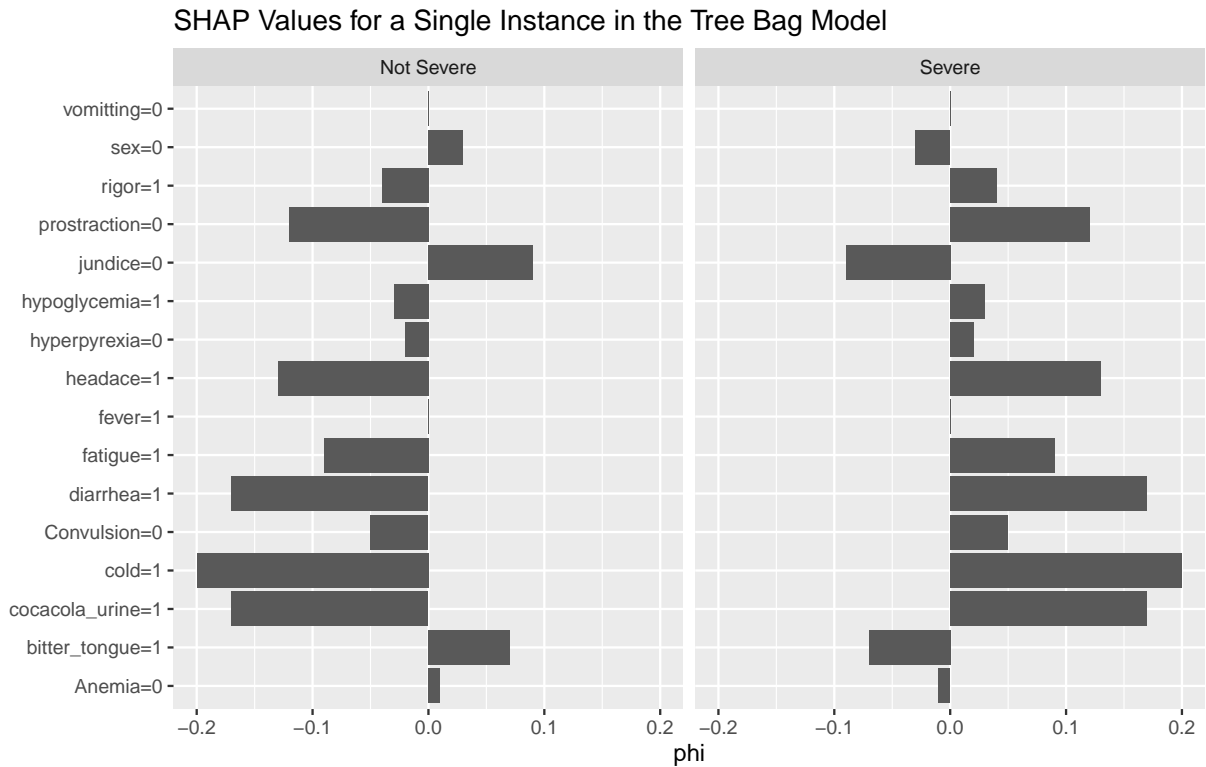
```
x_interest <- test[1, -which(names(test) == "severe_malaria")]
```

Compute SHAP values for the specific instance

```
shapleyTbag <- Shapley$new(predictorTbag, x_interest = x_interest)
```

Plot the SHAP values for this instance

```
shapleyTbag$plot() +  
  ggtitle("SHAP Values for a Single Instance in the Tree Bag Model")
```



Leftward (negative): Indicates the feature is pushing the model prediction towards a negative class (e.g., “Not Severe”). Larger absolute SHAP values mean a feature has a stronger influence on the prediction. Smaller SHAP values (close to zero) indicate that a feature has minimal influence on the model’s output for that instance